Hack-Nation

In collaboration with MIT Club of Northern California and MIT Club of Germany

CHALLENGE

04

The AI Scientist

From hypothesis to runnable experiment plan

POWERED BY

FULCRUM

FULCRUM SCIENCE

The Challenge

Turning a scientific question into a runnable experiment takes weeks of manual work — designing the protocol, estimating costs, sourcing materials, staffing the team. It's not the ideas that slow science down; it's the operations.

Here's how it typically works: an organisation — a pharma company, a university, a research funder — sends a brief to a specialist lab or Contract Research Organisation (CRO). The lab's scientists scope the work: they figure out what steps to follow, what to buy, how long it will take, and how much it will cost. Then they send back a proposal. A senior scientist who's run a similar experiment before can do this in hours. One who hasn't may take days — and the quality difference is real. A plan with the wrong chemical concentration or an unrealistic timeline can send a lab down the wrong path for weeks.

Your challenge is to build an AI-powered tool that compresses this process: take a scientific question as input, and generate a complete, operationally realistic experiment plan that a real lab could pick up and execute. You don't need to be a scientist to build this — you need to know how to build tools that make experts faster.

The mRNA moment	What accelerated COVID-19 vaccine development wasn't a new idea — the mRNA hypothesis had existed for decades. It was the ability to move from hypothesis to executable experiment at unprecedented speed. That's the gap you're closing.

What You're Building

A focused, end-to-end application with three stages:

1. Input	2. Literature QC	3. Experiment Plan
Natural language scientific question	Has this exact protocol been done before?	Full operational plan — the core deliverable

Primary Challenge: Full Experiment Planning

This is where you can make your mark. Can you build the tool that a real PI would actually want to use? Can you generate an experiment plan so complete, so operationally grounded, that a lab could pick it up on Monday and start running it by Friday? Here's what that looks like in practice — a plan that includes:

- **Protocol** — step-by-step methodology grounded in real published protocols
- **Materials and supply chain** — specific reagents, catalog numbers, suppliers
- **Budget** — realistic cost estimate with line items

- **Timeline** — phased breakdown with dependencies
- **Validation approach** — how success or failure will be measured

Quality bar	Would a real scientist trust this plan enough to order the materials and start running it? That is the standard.

Literature Quality Control Step

Before generating the plan, run a quick check: has this experiment, or something very close to it, been done before?

Think of this as a plagiarism check, but for science. It's not a deep literature review — it's a fast signal that tells the scientist whether they're breaking new ground or building on existing work. The output should be simple:

- A novelty signal: "not found", "similar work exists", or "exact match found"
- 1–3 relevant references where applicable, so the scientist can follow up

Use any literature source you like — arXiv, Semantic Scholar, protocol repositories, or others. The source matters less than the signal being accurate and fast.

The UI

Build a polished end-to-end interface. Submitted projects will be evaluated on both the quality of the generated plan and the quality of the experience.

The UI should make it easy to:

- Enter a scientific question in plain language
- See the literature QC result clearly before the plan is generated
- Read and navigate the full experiment plan
- Understand the budget, timeline, and materials at a glance

Stretch Goal: Scientist Review — Closing the Learning Loop

This is the hardest challenge in the brief, and the one with the highest ceiling.

The core idea: every time a scientist reviews and corrects a generated plan, that feedback becomes a training signal. Over time, the system gets better at generating plans for similar experiment types — learning from expert knowledge rather than just retrieving it.

What this looks like in practice:

- A structured review interface where a scientist can rate, correct, and annotate sections of the plan — protocol steps, reagent choices, budget lines, timeline assumptions
- A feedback store that captures corrections in structured form, tagged by experiment type and domain
- A generation layer that incorporates prior feedback when producing new plans of a similar type — at minimum as few-shot examples, at best as a lightweight fine-tuning loop

The demo that wins this stretch goal is one where a judge can watch the system generate a plan, a scientist leave structured corrections, and the next plan for a similar experiment visibly reflect those corrections — without being explicitly re-prompted.

Why this matters	A system that learns from scientist feedback compounds in value over time. Every review makes the next plan better. This is the difference between a tool and a platform.

Sample Inputs

Not sure what to test your tool with? Here are four example inputs — each is a real scientific hypothesis written the way a researcher would actually phrase it. Your tool should be able to take any of these and produce a full experiment plan.

Don't worry if the scientific terminology is unfamiliar — we've added plain-English translations for each one.

.

Diagnostics	A paper-based electrochemical biosensor functionalized with anti-CRP antibodies will detect C-reactive protein in whole blood at concentrations below 0.5 mg/L within 10 minutes, matching laboratory ELISA sensitivity without requiring sample preprocessing. *In plain English: can we build a cheap, fast blood test for inflammation that works without lab equipment?*
Gut Health	Supplementing C57BL/6 mice with Lactobacillus rhamnosus GG for 4 weeks will reduce intestinal permeability by at least 30% compared to controls, measured by FITC-dextran assay, due to upregulation of tight junction proteins claudin-1 and occludin. *In plain English: does a specific probiotic measurably strengthen the gut lining in mice?*
Cell Biology	Replacing sucrose with trehalose as a cryoprotectant in the freezing medium will increase post-thaw viability of HeLa cells by at least 15 percentage points compared to the standard DMSO protocol, due to trehalose's superior membrane stabilization at low temperatures. *In plain English: can we keep more cells alive when freezing them by swapping one preservative for another?*
Climate	Introducing Sporomusa ovata into a bioelectrochemical system at a cathode potential of −400mV vs SHE will fix CO_2 into acetate at a rate of at least 150 mmol/L/day, outperforming current biocatalytic carbon capture benchmarks by at least 20%. *In plain English: can a specific microbe be used to convert CO_2 into a useful chemical compound more efficiently than current methods?*

What makes these strong inputs	Each hypothesis names a specific intervention, states a measurable outcome with a threshold, gives a mechanistic reason, and implies a clear control condition. "AI will improve drug discovery" is a goal, not a hypothesis.

Hints and Resources

Protocol Repositories

- protocols.io: largest active repository, structured format: protocols.io
- Bio-protocol: peer-reviewed, linked to papers: bio-protocol.org
- Nature Protocols: premium detail: nature.com/nprot
- JOVE: video protocols with written transcripts: jove.com
- OpenWetWare: community protocols: openwetware.org

Supplier References

- Thermo Fisher application notes: thermofisher.com/us/en/home/technical-resources/application-notes.html
- Sigma-Aldrich technical bulletins: sigmaaldrich.com/US/en/technical-documents
- Promega protocols: promega.com/resources/protocols
- Qiagen protocols: qiagen.com/us/resources/resourcedetail?id=protocols
- IDT primer design and qPCR: idtdna.com/pages/tools

Reagent and Cell Line References

- ATCC cell line protocols: atcc.org
- Addgene cloning and transfection: addgene.org/protocols

Scientific Standards

- MIQE Guidelines for qPCR: ncbi.nlm.nih.gov/pmc/articles/PMC2737408

Contact: arun@fulcrum.science / jonas@fulcrum.science

What Good Looks Like

A participant enters: “Can we improve solar cell efficiency by testing alternative materials?”

The tool checks whether this protocol exists, surfaces 2 prior papers, flags it as “similar work exists”, then generates a full plan: a 6-step synthesis protocol grounded in protocols.io, a materials list with catalog numbers and a £12,000 budget estimate, and a 10-week timeline.

That is the bar. Build something a real scientist would want to use.